

Chapter (6-7)-2: Advanced Classification Methods: Extra

Richard Liu

May 3, 2020



Contents

- 1 K-D tree: Basic Examples and Illustration
- 2 SVM: An Application of KKT Conditions
- 3 Adaboost: Mathematical Principles
- 4 Bagging versus Boosting: Other Properties



Source

- Zhihua Zhou, Machine Learning
- CS61B, University of California, Berkeley



Section 1

K-D tree: Basic Examples and Illustration



KNN Graph

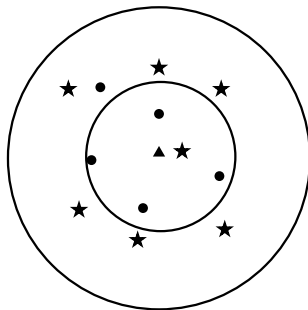


Figure: Graph of KNN



K-D Tree Graph

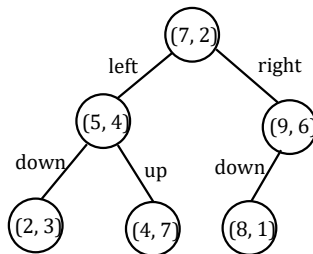
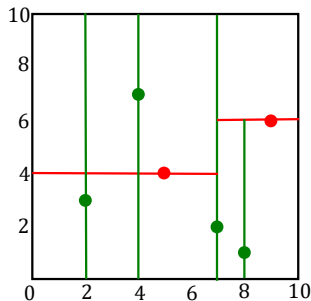


Figure: An illustration of K-D tree

- It divides the plane into different parts, first left/right then up/down, then left/right, ...



Other details

- Different inserting orders will lead to different trees. Sometimes the tree needs to be balanced.
- In some cases, the trees could be pruned.
- Compared with ordinary methods (just enumerate with time complexity $\mathcal{O}(n)$, K-D tree could provide a better performance if $k \ll n$ (Why?).
- **Python Code:** Here



Section 2

SVM: An Application of KKT Conditions



Mathematical Foundation of SVM Proof

Suppose that the general nonlinear programming problem is

$$\min_x f(x)$$

$$c_i(x) = 0, i \in \mathcal{E}, c_i(x) \leq 0, i \in \mathcal{I}$$



Mathematical Foundation of SVM Proof

Theorem 1: Karush-Kuhn-Tucker Conditions

Suppose that x^* is a local minimizer, that the function f and c_i are continuously differentiable, and that x^* is a regular point. Then there is a Lagrange multiplier vector λ^* with components $\lambda_i^*, i \in \mathcal{E} \cup \mathcal{I}$ such that the following conditions are satisfied

$$\nabla_x \mathcal{L}(x^*, \lambda^*) = 0$$

$$c_i(x^*) = 0, i \in \mathcal{E}, c_i(x^*) \geq 0, i \in \mathcal{I}$$

$$\lambda_i^* \geq 0, i \in \mathcal{I}, \lambda_i^* c_i(x^*) = 0, i \in \mathcal{I}$$



Section 3

Adaboost: Mathematical Principles



Pseudo-Code of Adaboost

输入: 训练集 $D = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\}$;
基学习算法 \mathcal{L} ;
训练轮数 T .

过程:

- 1: $\mathcal{D}_1(\mathbf{x}) = 1/m$.
- 2: **for** $t = 1, 2, \dots, T$ **do**
- 3: $h_t = \mathcal{L}(D, \mathcal{D}_t)$;
- 4: $\epsilon_t = P_{\mathbf{x} \sim \mathcal{D}_t}(h_t(\mathbf{x}) \neq f(\mathbf{x}))$;
- 5: **if** $\epsilon_t > 0.5$ **then break**
- 6: $\alpha_t = \frac{1}{2} \ln \left(\frac{1-\epsilon_t}{\epsilon_t} \right)$;
- 7:
$$\mathcal{D}_{t+1}(\mathbf{x}) = \frac{\mathcal{D}_t(\mathbf{x})}{Z_t} \times \begin{cases} \exp(-\alpha_t), & \text{if } h_t(\mathbf{x}) = f(\mathbf{x}) \\ \exp(\alpha_t), & \text{if } h_t(\mathbf{x}) \neq f(\mathbf{x}) \end{cases}$$
$$= \frac{\mathcal{D}_t(\mathbf{x}) \exp(-\alpha_t f(\mathbf{x}) h_t(\mathbf{x}))}{Z_t}$$
- 8: **end for**

输出: $H(\mathbf{x}) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(\mathbf{x}) \right)$

Figure: Adaboost Algorithm



Problem 1: Why such classifier?

- Consider a binary classification problem with $f(\mathbf{x})$ the true answer and $y_i = \{1, -1\}$.
- Note that we could write the whole classifier as

$$H(\mathbf{x}) = \sum_{t=1}^T \alpha_t h_t(\mathbf{x})$$

where $\{h_t(x)\}$, $t = 1, 2, \dots, T$ are base classifiers. Then we could optimize exponential loss function

$$l_{\text{exp}}(H|\mathcal{D}) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[e^{-f(\mathbf{x})H(\mathbf{x})}]$$

Explanation

It is easy to see $f(\mathbf{x})H(\mathbf{x})$ will be maximized when $H(\mathbf{x})$ achieves 100% accuracy.



Problem 1: Why such classifier?

Take derivatives w.r.t. $H(\mathbf{x})$ yields

$$\frac{\partial \ell_{\text{exp}}(H|\mathcal{D})}{\partial H(\mathbf{x})} = -e^{-H(\mathbf{x})}P(f(\mathbf{x}) = 1|\mathbf{x}) + e^{H(\mathbf{x})}P(f(\mathbf{x}) = -1|\mathbf{x})$$

F.O.C could get

$$H(\mathbf{x}) = \frac{1}{2} \ln \frac{P(f(\mathbf{x}) = 1|\mathbf{x})}{P(f(\mathbf{x}) = -1|\mathbf{x})}$$

Explanation

What is the definition of expectation with discrete cases?



Problem 1: Why such classifier?

That is to say

$$\begin{aligned}\text{sign}(H(\mathbf{x})) &= \text{sign} \left(\frac{1}{2} \ln \frac{P(f(\mathbf{x}) = 1|\mathbf{x})}{P(f(\mathbf{x}) = -1|\mathbf{x})} \right) \\ &= \arg \max_{y \in \{-1, 1\}} P(f(\mathbf{x}) = y|\mathbf{x})\end{aligned}$$

This means it achieves Bayesian optimal error rate, which means any other models could not lead to better results. In other words, we need to focus on $H(\mathbf{x})$.



Problem 2: How to choose α_i ?

In Adaboost, we use greedy algorithm, which means we only consider the t -th base classifier. In fact, we could write the formula as

$$\begin{aligned}\ell_{\text{exp}}(\alpha_t h_t | \mathcal{D}_t) &= \mathbb{E}_{\mathbf{x} \sim \mathcal{D}_t} [e^{-f(\mathbf{x})\alpha_t h_t(\mathbf{x})}] \\ &= \mathbb{E}_{\mathbf{x} \sim \mathcal{D}_t} [e^{-\alpha_t} \mathbb{I}(f(\mathbf{x}) = h_t(\mathbf{x})) + e^{\alpha_t} \mathbb{I}(f(\mathbf{x}) \neq h_t(\mathbf{x}))] \\ &= e^{-\alpha_t} P_{\mathbf{x} \sim \mathcal{D}_t}(f(\mathbf{x}) = h_t(\mathbf{x})) + e^{\alpha_t} P_{\mathbf{x} \sim \mathcal{D}_t}(f(\mathbf{x}) \neq h_t(\mathbf{x})) \\ &= e^{-\alpha_t} (1 - \epsilon_t) + e^{\alpha_t} \epsilon_t\end{aligned}$$

Take derivative of α_i could get the final result

$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$



Problem 3: How to choose right \mathcal{D} ?

- Ideally, we want every distribution to minimize the loss function. In this case, we want to minimize $\ell_{\text{exp}}(H_{t-1} + h_t | \mathcal{D})$ at time t , so we have

$$\begin{aligned}\ell_{\text{exp}}(H_{t-1} + h_t | \mathcal{D}) &= \mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \left[e^{-f(\mathbf{x})(H_{t-1}(\mathbf{x}) + h_t(\mathbf{x}))} \right] \\ &= \mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \left[e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})} e^{-f(\mathbf{x})h_t(\mathbf{x})} \right] \\ &\simeq \mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \left[e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})} \left(1 - f(\mathbf{x})h_t(\mathbf{x}) + \frac{f^2(\mathbf{x})h_t^2(\mathbf{x})}{2} \right) \right] \\ &= \mathbb{E}_{\mathbf{x} \sim \mathcal{D}} \left[e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})} \left(1 - f(\mathbf{x})h_t(\mathbf{x}) + \frac{1}{2} \right) \right]\end{aligned}$$

Explanation

Note that $f^2(\mathbf{x}) = h_t^2(\mathbf{x}) = 1$.



Problem 3: How to choose right \mathcal{D} ?

- Because we only consider the time t , the objective goal is to maximize

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})} f(\mathbf{x})h(\mathbf{x})]$$

- Consider (Why?)

$$\mathcal{D}_t(\mathbf{x}) = \frac{\mathcal{D}(\mathbf{x})e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})}}{\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})}]}$$

we will get

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})} f(\mathbf{x})h(\mathbf{x})] = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}_t} [f(\mathbf{x})h(\mathbf{x})]$$



Problem 3: How to choose right \mathcal{D} ?

For we have

$$f(\mathbf{x})h(\mathbf{x}) = 1 - 2\mathbb{I}(f(\mathbf{x}) \neq h(\mathbf{x}))$$

So the optimal base classifier at time t is

$$h_t(\mathbf{x}) = \arg \min_h \mathbb{E}_{\mathbf{x} \sim \mathcal{D}_t} [\mathbb{I}(f(\mathbf{x}) \neq h(\mathbf{x}))]$$

This means we could get the optimal solution with that



Problem 3: How to choose right \mathcal{D} ?

At last, by definition of \mathcal{D}_t , we have

$$\begin{aligned}\mathcal{D}_{t+1}(\mathbf{x}) &= \frac{\mathcal{D}(\mathbf{x})e^{-f(\mathbf{x})H_t(\mathbf{x})}}{\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [e^{-f(\mathbf{x})H_t(\mathbf{x})}]} \\ &= \frac{\mathcal{D}(\mathbf{x})e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})}e^{-f(\mathbf{x})\alpha_t h_t(\mathbf{x})}}{\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [e^{-f(\mathbf{x})H_t(\mathbf{x})}]} \\ &= \mathcal{D}_t(\mathbf{x}) \cdot e^{-f(\mathbf{x})\alpha_t h_t(\mathbf{x})} \frac{\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [e^{-f(\mathbf{x})H_{t-1}(\mathbf{x})}]}{\mathbb{E}_{\mathbf{x} \sim \mathcal{D}} [e^{-f(\mathbf{x})H_t(\mathbf{x})}]}\end{aligned}$$

So we get the updated formula.



Section 4

Bagging versus Boosting: Other Properties



- The goal of bagging is to make each base classifier achieve a better performance because the data could be used in more than one classifier.
- Bagging: Lower Variance. Boosting: Lower Bias.
- Random Forest is an extension of bagging, exploiting turbulence of features, so it usually behaves better than bagging.



Thank you!

